

ARCHIE BENN

HOW TO TRAIN YOUR MODEL

ASSESSING THE IMPACTS OF FOREST AGE ON THE SPATIAL
GENERALISATION OF RANDOM FOREST EVAPOTRANSPIRATION
MODELS

MSc Bioinformatics

Supervisor: Martin De Kauwe

May - September 2026



University of
BRISTOL

A dissertation submitted to the University of Bristol in accordance with the requirements of the degree of Master of Science by advanced study in Bioinformatics in the Faculty of Health and Life Sciences.

Abstract

Background

Accurate prediction of Evapotranspiration (ET) in regions using methods requiring minimal data input would enable estimations to be made in areas which are underrepresented globally due to the high costs involved in direct measurements. These estimations could, in turn, aid in more efficient use of freshwater and in predicting natural phenomena such as droughts. The development of Machine Learning (ML) models to predict ET has gained popularity in recent years, and being able to understand what links exist between a training dataset and a model which performs well on out-of-sample data should aid in the formation of better-predicting models, thus any explanatory variables which may aid in ET estimations should be explored. One under explored feature in ML models is forest age, which may impact rates due to age-related changes to canopy complexity and hydraulic conductance.

Methods

The present study considers the use of forest age as a feature in ML predictions of ET at forest sites contained within the FLUXNET by using methods to explore spatial generalisation of models trained with and without age included, alongside temporal generalisation to assess whether forest age may help predictions across time at single sites.

Results

Results from spatial generalisation suggest nuanced results, with age benefiting predictions at some held out test sites, and hurting at others. Further investigation into two sites with opposing results pointed towards x being a factor (MAYBE?). Temporal generalisation suggested forest age x y z .

Conclusion

These findings suggest that forest age may be useful to consider in ML ET modelling in the future, however considerations relating to tree species and the impacts of age on

tree dynamics within each species may be useful in determining whether to use forest age or not in predictions. Suggestions for future work include accurate age surveys to be carried out at other FLUXNET sites which may permit 'per-species models' to be formed and tested for use with forest age.

Key words: Machine Learning, Evapotranspiration, FLUXNET, Forest Age.

Word count total: 4186

Introduction: 231

Results: 1728

Discussion: 2226

Conclusion: 1

Dedication and Acknowledgments

I would like to thank my supervisor Martin de Kauwe for his help and guidance throughout the duration of this project.

Declaration

I declare that the work in this dissertation was carried out in accordance with the requirements of the University's Regulations and Code of Practice for Taught Programmes and that it has not been submitted for any other academic award. Except where indicated by specific reference in the text, this work is my own work. Work done in collaboration with, or with the assistance of others, is indicated as such. I have identified all material in this dissertation which is not my own work through appropriate referencing and acknowledgement. Where I have quoted or otherwise incorporated material which is the work of others, I have included the source in the references. Any views expressed in the dissertation, other than referenced material, are those of the author.

Archie Benn

September 3, 2026

Contents

Abbreviations	6
1 Aims and Objectives	7
2 Results	7
2.1 Flux Tower Site Clustering	7
2.2 Model Sensitivity	8
2.3 Train-Test Distance Analysis Using Sampled Training	8
2.4 Case Study of US-UMd and DE-Lnf	9
2.4.1 Observed vs. Predicted	9
2.4.2 Time Series Plots and SHAP Feature Importance	10
2.4.3 SHAP heatmap and Δ Predicted ET	10
2.4.4 Manipulating Training Ages	11
2.5 Results Summary	11
3 Discussion	11
3.0.1 Clustering	12
3.0.2 Model Sensitivity and Feature Selection	12
3.0.3 Model Generalisation	12
3.0.4 Forest Site Age as a Feature	13
3.0.5 Case Study	14
4 Conclusion	16
References	17
Figures and Tables	22
s1 Supplementary Materials and Methods	36
s1.1 Data Acquisition and Variables Setup	36
s1.2 Cluster Analysis	36
s1.3 Random Forest Modelling and Evaluation Metrics	37
s1.3.1 Model Sensitivity Testing	38
s1.3.2 Random Sampling of Training Datasets	38
s1.3.3 Two Site Case Study: US-UMd and DE-Lnf	39
s1.4 Site Metadata	39

Abbreviations

R^2 Coefficient of determination

DL Deep Learning

ET Evapotranspiration

LAI Leaf Area Index

LOGO Leave One Group Out

MAE Mean Absolute Error

ML Machine Learning

NEP Net Ecosystem Productivity

PAM Partitioning Around Medoids

RF Random Forest

RMSE Root Mean Square Error

SHAP SHapley Additive exPlanations

1 Aims and Objectives

The overarching aim of this research project is to investigate how the age of a forest may impact predictions of ET using ML in forest ecosystems. The stand age of a forest has been shown to impact ET dynamics [1], and therefore could alter how ML models form predictions of ET if age is included as a feature during training and inference.

Given the importance of predicting ET [2, 3] and the lack of infrastructure to measure ET directly with eddy covariance (EC) methods in many ecoregions [4], the first part of this project will focus on spatial generalisation, asking: *“If a model is built using readily available data from one set of sites, how well can it form predictions of ET rates at a geographically unseen site, and how does utilising forest age affect predictions?”*. To address this question, FLUXNET sites with forests aged by Besnard et al. [5] are used alongside a framework developed to quantify site similarities based on characteristics. Random Forest (RF) models were trained with and without forest age as a feature, and predictions of ET rates from both directly compared at unseen sites with ground-truth validation provided by EC measurements.

The objectives of this project are twofold: (1) To assess the use of forest age in models for spatial predictions of ET, and (2) to outline future directions for research in utilising forest age in ET predictions.

2 Results

2.1 Flux Tower Site Clustering

Cluster analysis was carried out in order to classify sites into groups based on their characteristics. This classification permits the predictability of ET, both within these characterised groups and across them, to be investigated. Gower’s distances, a measure between pairs of data points, from 0 = completely dissimilar to 1 = completely similar [6] (S1.2), were computed using the features listed in Table S1 and saved as a matrix. This matrix was passed to the Partitioning Around Medoids (PAM) algorithm [7] with the silhouette method to form clusters and identify medoid sites - those which have minimal dissimilarity to others within each cluster and thus can be seen as the representative sites of each cluster [8].

The assignment solution from PAM with $k = 6$ clusters can be seen in a cluster plot of the first two principal components (Figure 1a). The geographic distribution of clusters appears to follow an expected pattern with regards to climatic groups; for example, sites at high latitudes tended to be grouped in cluster 3, with Mediterranean and coastal sites typically grouped in cluster 5 (Figure 1b). This expected pattern of clusters suggests that the PAM clustering solution captures climate-relevant patterns to some degree, although

this was not formally tested. The six medoid sites were geographically distributed with four sites in North America and two sites in Europe (Figure 1b), and a range of site ages from 20 years (US-Me3) to 117 years (DE-Lnf).

Site characteristics ≥ 1 sd from the all site mean were considered as 'high/low' (≥ 1.5 sd as 'very high/very low') in order to qualitatively classify cluster characteristics, visualised in (Figure 2) and summarised in Table 1.

2.2 Model Sensitivity

Results from model sensitivity testing indicate mixed model performance across the medoid sites, based on Coefficient of determination (R^2), Root Mean Square Error (RMSE), and Mean Absolute Error (MAE) between the predicted and observed daily ET (Figure 3). Here, features were sequentially added to RF models, and these were re-trained with each additional feature before testing on all unseen medoid sites (S1.3.1).

Predictions of daily ET formed for CA-Obs, US-UMd, and DE-Lnf showed strong performance across the range of input features, while CA-TP3, IT-SRo, and US-Me3 was less predictable by the models overall.

CA-Obs and DE-Lnf showed very similar patterns of performance, with predictions at both sites typically strong and climbing across metrics up until $\cos_DOY$ was added (Peak CA-Obs: $R^2 = 0.90$; Peak DE-Lnf: $R^2 = 0.90$). Predictions at both sites then declined with the inclusion of forest age to the model (ΔR^2 with $Site_age$: CA-Obs = -0.07; DE-Lnf = -0.04). US-UMd deviated from this pattern, with performance improving further when $Site_age$ was included in the model (Peak US-UMd: $R^2 = 0.90$, ΔR^2 with $Site_age$ = +0.03). At all three of these sites, using only a single feature in the RF model, T_air , demonstrated strong to moderate predictive performance (CA-Obs: $R^2 = 0.76$; US-UMd: $R^2 = 0.64$; DE-Lnf: $R^2 = 0.63$).

Overall predictive performance at IT-SRo was poor (Peak IT-SRo: $R^2 = 0.28$). Performance at the remaining two sites, CA-TP3 and US-Me3, was particularly mixed and relatively poor (Peak CA-TP3: $R^2 = 0.55$; Peak US-Me3: $R^2 = 0.00$). R^2 fluctuated between predictions made with new features at these two sites, with performance at US-Me3 indicating that the model has no explanatory power at this site from a peak R^2 of 0.0.

Ratios of RMSE:MAE from peak predictive models at each medoid all fall between 1.33 and 1.48, suggesting that the impact of outlier ET predictions, and overall error distribution, was largely consistent from the best performing models at each site during this analysis.

2.3 Train-Test Distance Analysis Using Sampled Training

Figure 4 shows the paired results between the 'Mod' and the 'Mod + Age' models (described by Equation S1 and Equation S2 respectively) during the shuffled and sampled

training set analysis (S1.3.2). Including age had a significant impact across a range of training set compositions: all six medoid sites displayed significant differences in R^2 between 'Mod' and 'Mod + Age' predictions for daily ET (Wilcoxon signed-rank tests, $p < 0.01$ for all).

Mean Gower's and age distances between training set and test sites were calculated for each of the 100 runs to assess how training set composition influenced predictive performance (S1.3.2). Pooled site results follow an expected pattern: predictive accuracy falls as ecological and age dissimilarity increases between the training data and test sites in both models (Figure 5). For the 'Mod' model, the multiple linear regression was statistically significant ($R^2 = 0.60$, $F(2, 597) = 453.8$, $p < 0.001$), and both predictors independently explained R^2_{Mod} (Gower's: $\beta = -6.19$, $p < 0.001$; Age: $\beta = -0.0064$, $p < 0.001$). The 'Mod + Age' model was also significant ($R^2 = 0.54$, $F(2, 597) = 349$, $p < 0.001$), and both variables independently explaining $R^2_{Mod+Age}$ (Gower's: $\beta = -7.03$, $p < 0.001$; Age: $\beta = -0.0056$, $p < 0.001$). As site age was included as a variable during cluster analysis (Table S1), mean Gower's and age distances are not fully independent. However, very low multicollinearity ($VIF = 1.52$ for both predictors in each model) suggests each predictor contributes distinct information.

This similarity-performance relationship holds with all sites pooled, but is not present within each individual site. Simple linear models within sites suggest Gower's distance showed no significant relationship with R^2 at CA-TP3 in either model, and was significantly positively associated with R^2 in 'Mod + Age' predictions at US-Me3. Similarly, mean age distance was only significantly associated with R^2 at US-UMd for both models, with no significant relationship present at the other five medoid sites in either model.

2.4 Case Study of US-UMd and DE-Lnf

US-UMd and DE-Lnf share similar site qualities (Table S4) and both show high predictability for daily ET using ML models, yet the predictive performance responses in both section 2.2 and section 2.3 were opposite when age was added as a feature. Further analyses focused on these two sites (S1.3.3).

2.4.1 Observed vs. Predicted

Figure 6a and Figure 6b display scatter plots comparing the predicted and observed daily ET from each model at the sites. At DE-Lnf, the addition of age appears to cause low observed ET to be over-predicted, particularly in winter and autumn, in an opposite response to predictions of low observed ET at US-UMd. All models appear to under-predict peak daily ET in summer. Performance metrics followed the same pattern as in 2.2 and 2.3, with R^2 improving in US-UMd predictions ($R^2_{Mod} = 0.86$; $R^2_{Mod+Age} = 0.90$) and dropping

in DE-Lnf predictions ($R^2_{Mod} = 0.90$; $R^2_{Mod+Age} = 0.86$) with the addition of forest age as a feature.

2.4.2 Time Series Plots and SHAP Feature Importance

Time series plots of predictions from ‘Mod + Age’ (Figure 7a and Figure 7b) indicate strong cyclical predictive performance overall at both sites. Both models captured abnormal lows during summer months; for example, an ET rate which is typically seen during mid-winter was largely captured at US-UMd in summer 2011, demonstrating flexibility in these RF models. Despite this, extreme ET peaks were missed in consecutive years at US-UMd. For instance, a peak in summer 2010 was under-predicted by the model by nearly 2mm/day. Winters appeared to be problematic for the models, with ET estimations being consistently over-predicted and troughs missed in all years at both sites. Mean SHapley Additive exPlanations (SHAP) importance order was similar at both sites, with SW_rad and cos_D0Y dominating predictions in both models, followed by T_air and LAI. Site age was ranked as the 6th most important predictor at US-UMd and 5th at DE-Lnf. Mean SHAP values also indicate that site age had opposite impacts on predictions at both sites: age typically reduced predictions at US-UMd and increased them at DE-Lnf.

2.4.3 SHAP heatmap and Δ Predicted ET

To assess how SHAP values varied over time in the ‘Mod + Age’ model, mean weekly SHAP importance for features was plotted as a time-series heatmap, presented alongside a graph representing how ET predictions changed (Δ Predicted ET) between the ‘Mod’ and ‘Mod + Age’ models at both sites in Figure 8 (S1.3.3). Heatmaps suggest clear, yearly cyclical patterns of SHAP importance for all features apart from P_sum_14D and W_speed. Heatmap SHAP values for Site_age appear to line up with the Δ Predicted ET plot shifts from the $y = 0$ baseline for both sites. For example, Δ Predicted ET drops from the baseline most years at US-UMd appear to line up consistently with lower SHAP values for Site_age (Figure 8a). Similarly, at DE-Lnf, high peaks above the baseline correspond with high SHAP values for Site_age, such as in late 2007, mid 2010, and early 2012 (Figure 8b).

Scatter plots of Δ Predicted ET against Site_age SHAP values for the sites (Figure 9 and Figure 10) both indicate a strong, near linear relationship. Multiple linear regressions using SW_rad, cos_D0Y, and Site_age SHAP values as predictors confirmed that Site_age SHAP was the dominant predictor of Δ Predicted ET at both sites (US-UMd: $\beta = 2.47$, $SE = 0.049$, $p < 0.01$; DE-Lnf: $\beta = 2.49$, $SE = 0.036$, $p < 0.01$). In contrast, SW_rad and cos_D0Y SHAP values were minor, but significant predictors of Δ Predicted ET in all but SW_rad at US-UMd, which was not significant (full outputs in Table 2 and Table 3). This strongly suggests that prediction differences between ‘Mod’ and ‘Mod + Age’ were shaped by how Site_age is used within the RF models, rather than by the addition of this predictor altering decision splits in other features.

2.4.4 Manipulating Training Ages

Three new models were formed to investigate if real age-ET relationships were being learned by the models: (1) 'Mod' (2) 'Mod + Age' (3) 'Mod + Shuffled Age', where training site ages were randomly shuffled site-wise (S1.3.3).

Consistent with previous findings, daily ET predictive performance increased with the addition of age as a feature at US-UMd ($R^2_{Mod} = 0.86 \pm 0.00$; $R^2_{Mod+Age} = 0.90 \pm 0.00$), and decreased at DE-Lnf ($R^2_{Mod} = 0.90 \pm 0.00$; $R^2_{Mod+Age} = 0.86 \pm 0.00$) in models formed with all five NumPy seeds (Figure 12). At both sites, the shuffled age model moved predictive performance back towards the 'Mod' model (US-UMd $R^2_{Shuffled} = 0.85 \pm 0.03$; DE-Lnf $R^2_{Shuffled} = 0.89 \pm 0.02$).

2.5 Results Summary

Site age as a feature in RF models caused site-specific and non-uniform effects on daily ET predictions in all analyses. Train-test similarity analyses demonstrated that predictive accuracy declined as ecological and age dissimilarity increased, although these relationships were not as robust when assessed within individual sites. A case study of two sites identified that including age improved predictive performance at US-UMd but reduced it at DE-Lnf, despite both sites demonstrating high baseline predictability. Time-series SHAP analysis revealed a temporally consistent impact on ET predictions from site age at both sites, and modelling SHAP values against the change in ET predictions between 'Mod' and 'Mod + Age' strongly suggested site age was the dominant predictor driving these differences. Shuffling ages between sites during training caused predictive accuracy to shift from 'Mod + Age' back towards 'Mod' levels at both sites, despite age being present as a predictor. Together, these results show that the effect of site age as a feature in RF models is nuanced, site-specific, and detectable across numerous modelling approaches for predictions of daily ET.

3 Discussion

This study used ML models to form predictions of daily ET at FLUXNET forest sites across Europe and North America. While many previous studies have used FLUXNET with ML for ET predictions [empty citation], using the age of a forest as a feature in these models has not been explored previously, and this study represents the first use of this feature for predictive ET modelling using ML methods.

In all analyses contained within this study, site age contributed a meaningful impact on predictions of daily ET formed by RF models. The impact of using site age resulted in

nuanced predictive accuracy changes, with large variation in the direction and magnitude of these performance shifts between sites in most analyses.

3.0.1 Clustering

Cluster analysis with PAM has been used successfully in hydro-ecology [9, 10], but is rarely used in ML ET predictive modelling. In this study, model performance across clusters, likely because training data pooled sites from ecologically distant sites. Ferreira, Cunha, and Fernandes Filho [11] avoided this issue by training specific models and testing internally within each cluster. Clustering here was also defined by variables within the FAO56-PM equation [12], ensuring that the resulting clusters are relevant to ET estimations. While not formally verified, PAM clustering using Gower's distances in this study appeared to capture climate-relevant patterns, and may offer a useful framework to quantify regional diversity in future ET generalisation modelling.

3.0.2 Model Sensitivity and Feature Selection

Feature selection has direct implications in model accessibility, interpretability, and predictive accuracy [3]. Sensitivity testing at cluster medoids provided evidence that feature selection impacts performance differently across sites, supporting the idea that no universally optimal feature set exists for ET predictions [3]. Importantly, the addition of site age was never neutral in terms of performance, indicating that site age can alter model decision structures. Models trained using easily measured variables can improve accessibility in data sparse regions, but may limit accuracy compared to models built using richer data sources, with Nayak et al. [13] demonstrating an R^2 improvement from 0.78 to 0.98 when moving from a data-limited to data-rich Deep Learning (DL) models. In this study, predictable sites such as US-UMd, DE-Lnf, and CA-Obs could reach a high level of predictive accuracy using 2-3 features, in agreement with Nayak et al. [13]. Feature selection here relied on manual selection from variables understood to impact ET (Table S3), and were fixed across all analyses and models in this study which may have impacted performance inconsistently across test sites. A flexible approach with systematic selection methods, as suggested in Liyew et al. [3], can identify and select the best-performing features to reduce model complexity. Stapleton, Eichelmann, and Roantree [14] used this method to enhance model performance and reduce dimensionality, critically without requiring advanced domain knowledge. Applying forest age using this approach would allow more standardised assessments of performance across training datasets and reduce user-selection bias.

3.0.3 Model Generalisation

Predictive performance using sampled training datasets was highly variable across medoids. In contrast, Ahmadi et al. [15] reported far greater generalisation capacity, likely because

their models only considered sites within the Central Valley of California - a climatically homogeneous region, similar to the cluster-specific approach in Ferreira, Cunha, and Fernandes Filho [11], training and testing sites were confined to much smaller and similar geographic regions; all 55 sites from Ahmadi et al. [15] were contained within the Central Valley of California - a climatically homogeneous region. Restricting training and test data to ecologically similar sites likely enabled greater generalisation, whereas the present study used more heterogeneous training data. This aligns with regression analysis in section 2.3, where greater ecological dissimilarity between training sets and medoid sites was associated with reduced predictive accuracy. Although this relationship was not consistent within individual test sites, the pooled-site results suggested heterogeneity in training data can constrain model generalisation.

3.0.4 Forest Site Age as a Feature

This study was positioned by previous work carried out in Besnard et al. [5] where forest age was used as a predictor in Net Ecosystem Productivity (NEP) estimations using RF models. While current literature surrounding the use of forest age in ML ET models is sparse, this prior work was able to establish forest age as a dominant factor in NEP spatio-temporal variability at a global scale, providing evidence that models are able to learn from forest age for flux-predictions using RF.

The present study found more nuanced and site-specific results while investigating ET predictions, despite using 50 FLUXNET forest sites which were also used and dated in the NEP study. For feature selection, Besnard et al. [5] carried out a similar method (reversed) with features sequentially removed, and found that removing forest age always substantially decreased model performance, contrasting to the current study in which forest age appeared to have a less substantial impact. There are many possible explanations for the contrasting impact observed on the two fluxes in both of these studies. Crucially, Besnard et al. [5] predicted carbon fluxes rather than water fluxes. Carbon fluxes such as NEP across forest ecosystems may be more directly correlated with the age of those forests than any ET-age relationship, encouraging the RF models to place greater importance on forest age when forming NEP predictions. Furthermore, 'age' in the models used in Besnard et al. [5] was not solely raw age as in the present study. Instead, a pre-defined nonlinear relationship between NEP and age [16] was fitted, and a predicted value per site year from this was treated as a separate variable in the models alongside raw age. In contrast, the present study was designed to let the RF models draw any ET-age relationships from raw age alone. Oishi et al. [1] demonstrated an age-ET relationship but this was particular to a specific region and species composition; no generalisable equation form exists as with NEP in [16]. Even if one did exist, forcing it as a fixed input into a model, as in Besnard et al. [5], risks imposing model outputs which are constrained to a specific pattern and may not hold true across all sites. The use of raw age alone in

this study likely led to a noisier age-target relationship across models, and could explain the highly site-dependent impacts on predictions compared to the clean global signal in Besnard et al. [5], independent of whether the target was NEP or ET.

3.0.5 Case Study

US-UMd and DE-Lnf were highly predictable sites but had opposing reactions in predictive performance with the addition of age. At both sites, SHAP model importance at both sites was dominated by shortwave radiation, Leaf Area Index (LAI), and air temperature - consistent with prior ET modelling [3].

'Mod' and 'Mod + Age' predictions at both sites followed a consistent pattern: summer extremes were under-predicted and winter lows were over-predicted (Figure 6). SHAP values for `Site_age` were opposite in direction at the two sites, being positive at DE-Lnf and negative at US-UMd (Figure 7). This effect was strongly seasonal, with absolute SHAP values peaking during winter and tracking Δ Predicted ET over time (Figure 8). Δ Predicted ET values are caused by the addition of site age, but as RF models can adapt how decisions are split when new features are added, changes in predictions may be indirectly driven by features already included in 'Mod' rather than `Site_age` itself [empty citation]. Regressions of Δ Predicted ET against SHAP values for `Site_age`, `SW_rad`, and `cos_D0Y` ruled this out, and found `Site_age` to be highly significant and by far the strongest predictor at both sites (Table 2 and Table 3), confirming that the seasonal importance of `Site_age` on winter ET predictions - increasing them away from observed values at DE-Lnf and decreasing predictions towards them at US-UMd - drove the opposing shifts in R^2 . The influence of age in summer was comparatively small (Figure 8), possibly because of the greater presence of variable features in these months, and so it was insufficient to offset the winter-driven R^2 changes.

The exact mechanism for these opposing effects remains unclear, and the limited explainability of ML models makes this difficult to determine with confidence [17]. Forest age varies substantially between sites but little within each, unlike other dynamic predictors in the models. This raises the possibility of models using age less as a daily predictor and more as a site-level offset for ET, shifting predictions in a constant direction based on age relative to the training set. The SHAP dependence plot supports this: SHAP importance direction for `Site_age` is almost stepwise around the all-site mean forest age (Figure 11), implying that some age-ET relationship where ET is greater at older sites is being learned. US-UMd and DE-Lnf ages fall on opposite sides of this mean, explaining their divergent responses in R^2 when age is included. Random shuffling of ages between sites before carrying out a train-test cycle suggests that a relationship involving true age-ET was established within the models: performance shifted back towards 'Mod' levels, and the directional shift at the sites was lost despite the inclusion of `Site_age` in 'Mod + Shuffled Age'. Figure 13 shows a nonlinear relationship between forest age and both total

and 90-day peak ET sums across all site-years. This follows a similar overall shape as the SHAP dependence plot and could represent an age-ET relationship being learned by the models.

Why this relationship exists is difficult to establish. With a small number of heterogeneous sites spanning multiple tree species and climate zones, a clear age-ET relationship which fits all sites is unlikely to emerge, and models are likely trying to fit a noisy signal to unseen sites. Developmental dynamics will vary by species and climate [[empty citation](#)], and characteristics known to impact ET rates - such as stomatal conductance [[empty citation](#)] and canopy density [[empty citation](#)] - are likely to vary inconsistently with age in the sites used here. This suggests that the impact of forest age on ET should be more informative if tree species and climate zone are shared between training data and test sites, allowing potential age-ET relationships to emerge more clearly.

Address limitations and potential impact on results 300 words

This site heterogeneity reflects structural constraints in the model development used here. For models trained on sites spanning different tree species and climate zones, learning consistent relationships is inherently difficult if these are not true in the majority of training data [[empty citation](#)], even if these relationships exist in certain conditions. This is supported by the highly varied predictive accuracy across different clusters ([Figure 5](#)) - ultimately, underrepresented sites in the feature space during training were harder to predict. Uncertainty in the determination of site ages by Besnard et al. [[5](#)], FLUXNET accuracy issues [[18](#)], and the use of a fixed set of variables and hyperparameters during RF training (section [S1.3](#)) likely did not help with overall predictive accuracy and uncertainty, but the dominant constraint remains the ecological diversity of the training data. Although an age-ET relationship was learned here by the models and demonstrated clear impacts of age on predictions, this underlying relationship is likely messy and confounded by species and climate-specific dynamics. As a result, age-induced prediction shifts here may not be reflective of any true biophysical association. This necessitates reflecting on the reliance of R^2 as a measure for a 'good' model. R^2 shifts appeared to only result from season-specific predictions either driving ET estimates up or down based on where each test site is placed in the age-ET space of the training data. Neither 'Mod' nor 'Mod + Age' is necessarily better or worse in terms of understanding ET dynamics; the influence of forest age becomes difficult to interpret under such heterogeneity, irrespective of R^2 .

Implications and directions on/for future research 250 words

Developing ML models specific to tree species or cover types would reduce ecological confounding and allow cleaner age-ET relationships to emerge. The limited availability of

age-dated FLUXNET sites constrained the development of species-specific models here, but hybrid ML modelling, in which physics-based methods are paired with ML models, may offer a solution to this issue Liyew et al. [3]. Jiang et al. [19] used this framework to simulate ET values, allowing the target variable of RF models to be simulated in a data-deficient basin. This approach may increase uncertainty as the target is not directly measured, but it could significantly expand the training data spatially. If combined with improved forest age information through field-based surveys or LiDAR-based identification of species and tree turnover events [fischer2025global], the formation of species-specific models for ET predictions may be possible. These models would learn cleaner age-ET relationships, providing a stronger basis for evaluating whether forest age could be a useful predictor on a species-by-species basis.

A further direction for future work is to clarify whether age acts as an ‘explainer’ or ‘explained’ variable in ET prediction. This study aimed to identify if forest age may indirectly capture properties of forests relating to ET dynamics over their lifetime. If other measurements are able to capture these forest properties relating to ET dynamics directly however, the value in forest age becomes redundant as a feature and age can be thought of as ‘explained’. For example, it might be possible that canopy complexity over time can be fully captured through variables such as LAI or NDVI, bypassing use of age here entirely - particularly as these can be acquired more easily. Conversely, if forest age encodes information that is not captured by other measurements, or if it provides a more simple index by summarising multiple interacting biophysical processes, it can be thought of as an ‘explainer’ and may remain beneficial in modelling. Determining which role forest age plays should direct the necessity for further work. Ultimately, refining which features meaningfully improve ET predictions, including forest age, supports the wider goal of enabling accurate and low-cost ET prediction in data-sparse regions where direct measurement remains impractical.

4 Conclusion

References


- [1] AC Oishi, SO Denham, ST Brantley, KA Novick, PV Bolstad, and CF Miniati. “Quantifying the effects of stand age on components of forest evapotranspiration.” In: (2020).
- [2] Shima Amani and Hossein Shafizadeh-Moghadam. “A review of machine learning models and influential factors for estimating evapotranspiration using remote sensing and ground-based data”. In: *Agricultural water management* 284 (2023), p. 108324.
- [3] Chalachew Muluken Liyew, Stefano Ferraris, Elvira Di Nardo, and Rosa Meo. “A review of feature selection methods for actual evapotranspiration prediction”. In: *Artificial Intelligence Review* 58.10 (2025), p. 292.
- [4] Timothy Hill, Melanie Chocholek, and Robert Clement. “The case for increasing the statistical power of eddy covariance ecosystem studies: why, where and how?” In: *Global Change Biology* 23.6 (2017), pp. 2154–2165.
- [5] Simon Besnard, Nuno Carvalhais, M Altaf Arain, Andrew Black, Sytze De Bruin, Nina Buchmann, Alessandro Cescatti, Jiquan Chen, Jan GPW Clevers, Ankur R Desai, et al. “Quantifying the effect of forest age in annual net forest carbon balance”. In: *Environmental Research Letters* 13.12 (2018), p. 124018.
- [6] John C Gower. “A general coefficient of similarity and some of its properties”. In: *Biometrics* (1971), pp. 857–871.
- [7] Peter J Rousseeuw and L Kaufman. “Finding groups in data”. In: *Hoboken: Wiley Online Library* 1 (1990), p. 371.
- [8] Erich Schubert and Peter J Rousseeuw. “Fast and eager k-medoids clustering: O (k) runtime improvement of the PAM, CLARA, and CLARANS algorithms”. In: *Information Systems* 101 (2021), p. 101804.
- [9] Dele Hou, Yihui Qin, Ziqing Zhang, and Yongle Chen. “Site-specific growth-climate relationships of a widespread coniferous species across mountain areas”. In: *Forest Ecology and Management* 587 (2025), p. 122731.
- [10] Arthur Amaral e Silva, Leonardo Campos de Assis, Laura Coelho de Andrade, Juliana Ferreira Lorentz, Júlio César de Oliveira, Maria Lucia Calijuri, and Italo Oliveira Ferreira. “Multivariate Statistical Analysis of Rainfall Variability in Brazil: Assessing Climatic and Environmental Drivers of Precipitation”. In: *Remote Sensing Applications: Society and Environment* (2025), p. 101849.

- [11] Lucas Borges Ferreira, Fernando França da Cunha, and Elpídio Inácio Fernandes Filho. “Exploring machine learning and multi-task learning to estimate meteorological data and reference evapotranspiration across Brazil”. In: *Agricultural water management* 259 (2022), p. 107281.
- [12] Richard G Allen, William O Pruitt, James L Wright, Terry A Howell, Francesca Ventura, Richard Snyder, Daniel Itenfisu, Pasquale Steduto, Joaquin Berengena, Javier Baselga Yrisarry, et al. “A recommendation on standardized surface resistance for hourly calculation of reference ETo by the FAO56 Penman-Monteith method”. In: *Agricultural water management* 81.1-2 (2006), pp. 1–22.
- [13] Ajit Kumar Nayak, A Sarangi, S Pradhan, RK Panda, NM Jeepsa, BS Satpathy, and Mithlesh Kumar. “Estimation of daily reference evapotranspiration using machine learning and deep learning techniques with sparse meteorological data”. In: *Agricultural Research* (2026), pp. 1–14.
- [14] Adam Stapleton, Elke Eichelmann, and Mark Roantree. “A framework for constructing machine learning models with feature set optimisation for evapotranspiration partitioning”. In: *Applied Computing and Geosciences* 16 (2022), p. 100105.
- [15] Arman Ahmadi, Andre Daccache, Minxue He, Peyman Namadi, Alireza Ghaderi Bafti, Prabhjot Sandhu, Zhaojun Bai, Richard L Snyder, and Tariq Kadir. “Enhancing the accuracy and generalizability of reference evapotranspiration forecasting in California using deep global learning”. In: *Journal of Hydrology: Regional Studies* 59 (2025), p. 102339.
- [16] Brian D Amiro, Alan G Barr, JG Barr, T Andrew Black, R Bracho, M Brown, J Chen, KL Clark, KJ Davis, AR Desai, et al. “Ecosystem carbon dioxide fluxes after disturbance in forests of North America”. In: *Journal of Geophysical Research: Biogeosciences* 115.G4 (2010).
- [17] Vikas Hassija, Vinay Chamola, Atmesh Mahapatra, Abhinandan Singal, Divyansh Goel, Kaizhu Huang, Simone Scardapane, Indro Spinelli, Mufti Mahmud, and Amir Hussain. “Interpreting black-box models: a review on explainable artificial intelligence”. In: *Cognitive Computation* 16.1 (2024), pp. 45–74.
- [18] Timothy Hill, Melanie Chocholek, and Robert Clement. “The case for increasing the statistical power of eddy covariance ecosystem studies: why, where and how?” In: *Global Change Biology* 23.6 (2017), pp. 2154–2165.
- [19] Xiaoman Jiang, Guoqiang Wang, Yuntao Wang, Jiping Yao, Baolin Xue, and Yinglan A. “A hybrid framework for simulating actual evapotranspiration in data-deficient areas: A case study of the Inner Mongolia Section of the Yellow River Basin”. In: *Remote Sensing* 15.9 (2023), p. 2234.


- [20] Ahlmann-Eltze Constantin and Indrajeet Patil. “ggsignif: R Package for Displaying Significance Brackets for 'ggplot2'”. In: *PsyArxiv* (2021). doi: [10.31234/osf.io/7awm6](https://doi.org/10.31234/osf.io/7awm6). URL: <https://psyarxiv.com/7awm6>.
- [21] Marek Hlavac. *stargazer: Well-Formatted Regression and Summary Statistics Tables*. R package version 5.2.3. Social Policy Institute. Bratislava, Slovakia, 2022. URL: <https://CRAN.R-project.org/package=stargazer>.
- [22] Gilberto Pastorello, Carlo Trotta, Eleonora Canfora, Housen Chu, Danielle Christianson, You-Wei Cheah, Cristina Poindexter, Jiquan Chen, Abdelrahman Elbashandy, Marty Humphrey, et al. “The FLUXNET2015 dataset and the ONEFlux processing pipeline for eddy covariance data”. In: *Scientific Data* 7.1 (July 2020), p. 225. issn: 2052-4463. doi: [10.1038/s41597-020-0534-3](https://doi.org/10.1038/s41597-020-0534-3). URL: <https://doi.org/10.1038/s41597-020-0534-3>.
- [23] R Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing. Vienna, Austria, 2026. URL: <https://www.R-project.org/>.
- [24] Hadley Wickham, Mara Averick, Jennifer Bryan, Winston Chang, Lucy D'Agostino McGowan, Romain François, Garrett Golemund, Alex Hayes, Lionel Henry, Jim Hester, et al. “Welcome to the tidyverse”. In: *Journal of Open Source Software* 4.43 (2019), p. 1686. doi: [10.21105/joss.01686](https://doi.org/10.21105/joss.01686).
- [25] Juergen Knauer, Tarek Sebastian El-Madany, Soenke Zaehle, and Mirco Migliavacca. “Bigleaf - An R package for the calculation of physical and physiological ecosystem properties from eddy covariance data”. In: *PLoS ONE* 13.8 (2018), e0201114. doi: [10.1371/journal.pone.0201114](https://doi.org/10.1371/journal.pone.0201114). URL: <https://doi.org/10.1371/journal.pone.0201114>.
- [26] R Myneni, Y Knyazikhin, and T Park. *MCD15A3H MODIS/Terra+ Aqua Leaf Area Index/FPAR 4-day L4 Global 500m SIN Grid V006, NASA EOSDIS Land Processes DAAC [data set]*. 2015.
- [27] Koen Hufkens. “bluegreen-labs/MODISTools: MODISTools v1. 1.4 (CRAN release)”. In: *Zenodo* (2023).
- [28] Sabine Landau and I Chis Ster. “Cluster analysis: overview”. In: *International encyclopedia of education* 3 (2010), pp. 72–83.
- [29] Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, Peter Prettenhofer, Ron Weiss, Vincent Dubourg, et al. “Scikit-learn: Machine learning in Python”. In: *the Journal of machine Learning research* 12 (2011), pp. 2825–2830.

- [30] Takuya Akiba, Shotaro Sano, Toshihiko Yanase, Takeru Ohta, and Masanori Koyama. “Optuna: A next-generation hyperparameter optimization framework”. In: *Proceedings of the 25th ACM SIGKDD international conference on knowledge discovery & data mining*. 2019, pp. 2623–2631.
- [31] Davide Chicco, Matthijs J Warrens, and Giuseppe Jurman. “The coefficient of determination R-squared is more informative than SMAPE, MAE, MAPE, MSE and RMSE in regression analysis evaluation”. In: *Peerj computer science* 7 (2021), e623.

Figures and Tables



(a) Cluster plot indicating the site positions in the space of the first two principal components, with coloured clusters indicating the PAM solution. Axes indicate the % variance explained by each dimension.



(b) Map of all FLUXNET sites and associated clusters in this study, with cluster medoid site IDs highlighted. Visual inspection of cluster geographic distribution aligns with an expected pattern with regards to climatic groupings, suggesting the clustering solution captures climate-relevant patterns.

Figure 1: PAM solution from analysis at $k = 6$ clusters and using Gower's distances shown on a cluster plot in [a](#), and geographic cluster distributions shown in [b](#).




Figure 2: Heatmap of cluster medoid sites and scaled feature values. Scaled features ≥ 1.5 standard deviations from the mean of all sites were classed as 'very' (high or low, depending on direction), while features between 1 and 1.5 standard deviations from the mean were classed as high or low, depending on direction, in order to characterise clusters.

Table 1: Main characteristics of each cluster defined during cluster analysis

Cluster	Medoid Site	Defining Characteristics
1	DE-Lnf	Temperate sites without significant features.
2	US-UMd	Continental sites without significant features.
3	CA-Obs	Continental sites with low mean air temperatures.
4	CA-TP3	Continental sites with high precipitation.
5	IT-SRo	Temperate sites with high air temperatures.
6	US-Me3	Temperate sites with very high VPD, high SW radiation, low stand age, low wind speeds, and low precipitation.




Figure 3: R^2 , RMSE, and MAE for medoid cluster site predictions of daily ET during sensitivity testing (ordered left-right and top-bottom by peak R^2). Models were built and tested using a consecutively greater number of input features from FLUXNET and LAI data (for example, the ‘+VPD model’ incorporates Tair, SW_rad, and VPD as training and input features). R^2 performance at only half of the sites was moderate or strong, and suggests that adding features up to LAI improved prediction accuracy at four of the six medoid sites. Adding site age as a final feature led to mixed performance across ET predictions made at medoid sites, with only two of the six benefiting in terms of predictive performance from this input feature.




Figure 4: Paired box-plots indicate the impact on R^2 when age is added as a feature to RF models across 100 repeated runs of randomly sampled training datasets. Points are connected on a per-run basis where training dataset site composition is identical. Wilcoxon signed-rank tests suggest model performances differ significantly at all test sites, as indicated by the significance annotation on each facet (** signifies $p < 0.01$, *** signifies $p < 0.001$ [20]). DE-Lnf is shown large for readability as an example, with other cluster medoid sites shown at a reduced scale for comparison.






Figure 5: Plots displaying how R^2 of ET predictions relates to the mean Gower's distance and mean age distance from the training sites to each medoid test site with all sites together (100 runs). All fitted lines are a simple single linear regressions. The top plots are results from Mod, and the lower plots are results from Mod + Age, with results from all medoid test sites shown per plot. Visualisation indicate a correlation between R^2 and the mean Gower's distance in both Mod and Mod + Age, implying that models formed from datasets more closely related to the test sites produce more accurate predictions of ET. A relationship between R^2 and the mean age distance can be seen on the side plots, where predictions at test sites whose ages were further from the mean age of the training sites were typically lowered in both models.




(a) Subset of 5000 points comparing the observed daily ET and predicted daily ET from the two RF models with a 1:1 line shown in red (perfect model line). 'Mod + Age' improved prediction metrics overall compared to 'Mod', with a higher R^2 and lower RMSE and MAE. Visually, points appear tighter to the 1:1 line at both extremes, although extreme highs in observed daily ET rates are often under-predicted in both models.




(b) Subset of 2500

Figure 6: Case Study plots for site US-UMd.




(a) Time series observed daily ET overlaid with the predicted daily ET using the 'Mod + Age' RF model seen on the left, with average SHAP feature importance values across all predictions made with this model on the right. Time series plot suggests a strong fit to the observed data but with peak ET rates in extreme summers being under-predicted, in line with the shape seen in plot [a](#). Feature importance with SHAP...[CONTINUE]




(b) M.

Figure 7: Case Study plots for site DE-Lnf.



(a)



(b) yer

Figure 8: C




Figure 9: yh

Table 2: Multiple linear regression model results for US-UMd [21]

<i>Dependent variable:</i>	
	ET_diff
Site_age	2.467*** (0.049)
cos_DOY	-0.064*** (0.007)
SW_rad	0.004 (0.005)
Constant	0.131*** (0.004)
Observations	1,250
R ²	0.766
Adjusted R ²	0.766
Residual Std. Error	0.058 (df = 1246)
F Statistic	1,361.375*** (df = 3; 1246)

Note: *p<0.1; **p<0.05; ***p<0.01



Figure 10: bbb

Table 3: Multiple linear regression model results for DE-Lnf [21]

<i>Dependent variable:</i>	
	ET_diff
Site_age	2.494*** (0.036)
cos_DOY	-0.036*** (0.006)
SW_rad	0.043*** (0.007)
Constant	-0.117*** (0.003)
Observations	1,250
R ²	0.830
Adjusted R ²	0.829
Residual Std. Error	0.060 (df = 1246)
F Statistic	2,025.979*** (df = 3; 1246)

Note: *p<0.1; **p<0.05; ***p<0.01




Figure 11: SHAP dependence plot for `Site_age`. Test site SHAP values are highlighted in colour and represent model decisions for this feature on the unseen test sites when forming predictions for daily ET. Training site SHAP values are represented by the black points, and the red dashed line represents the mean forest age at all sites in this study. An almost stepwise, clear separation is visible for SHAP importance around the mean value for raw site age: a negative importance is typically placed at sites where forest age falls below the all-site mean, and positive at sites above this value. As the forest ages of US-UMd and DE-Lnf fall on separate sides of the all-site mean this may explain the model decisions in pushing ET predictions up (DE-Lnf) or down (US-UMd), with season-specific importance further relating to changes in overall R^2 .






Figure 12: Bar plot showing the effect of different model setups on R^2 at DE-Lnf and US-UMd, with error bars indicating standard error. To assess whether site age was acting as a proxy for other site-level characteristics in these models, a model which randomly shuffled ages between sites was also formed using identical seeds and LOGO approach. Predictions from the 'shuffled age' model resulted in an increased R^2 at DE-Lnf compared to the 'Mod + Age' model, where an opposite effect was observed at US-UMd, consistent with previous findings between the two models. This suggests that both models are learning from true age-ET relationships



(a) ET sum for site years plotted against the raw forest age at sites. A GAM with $k = 7$ is fitted as a visual aid.



(b) Sum 90D ET

Figure 13: A GAM with $k = 7$ is fitted as a visual aid.

S1 Supplementary Materials and Methods

All scripts used in this study, alongside a guide of how to re-run the full analysis, can be found at: <https://github.com/archiebenn/forest-age-ET.git>

S1.1 Data Acquisition and Variables Setup

Flux tower data were downloaded from the FLUXNET2015 dataset [22] and processed using R (version 4.5.3) [23] with the tidyverse [24] suite of packages. Flux data were filtered to only keep North American and European sites whose forest ages are dated in Besnard et al. [5], with rows of variables within these sites filtered based on their provided FLUXNET QC flag value. ET rates were calculated row-wise from the provided LE (latent heat flux) and T_{air} columns using the `LE.to.ET()` function from `bigleaf` [25] and multiplying this value by 86400 (s d^{-1}). Therefore, ET represents actual ET (AET) throughout this study. Four-day LAI data (NASA product: MCD15A3H [26]) were retrieved using site coordinates via `MODISTools` [27], filtered on QC, and linearly interpolated between the four-day periods of LAI to provide a continuous daily value, before joining by date and site to the main FLUXNET dataset.

Site stand ages retrieved from Besnard et al. [5] are static despite multi-year data ranges often contained at these sites in the FLUXNET. As such, site ages were extrapolated across the range of data at each site to act as a continuous metric, with the provided age set to the first row of the FLUXNET data for each site. Köppen climate sites were retrieved and attached based on site latitude and longitude, and these classifications were further simplified to their main groups (Temperate and Continental in this case). To represent the day of year (DOY), $\sin()$ and $\cos()$ functions were applied to the radian value of the day within the year ($\frac{2\pi\text{DOY}}{365}$) and together provide a cyclical representation for time throughout the year. A two week cumulative sum of precipitation prior to the day of measurement was also calculated and attached per row, with the first 13 days of data at each site then being discarded.

S1.2 Cluster Analysis

Cluster analysis is an exploratory technique which aims to classify objects into as-yet-unknown groups based on their similarities [28]. To group sites in this study using cluster analysis, a Gower's distance matrix was first formed. Gower's distances are a similarity measure which can handle both numerical and categorical data to provide a distance value between two data points (sites) from 0 (completely dissimilar) to 1 (completely similar) [6]. A table of the site features used for these distance calculations can be seen in [Table S1](#), and in this study Gower's distance is treated as a measure of ecological dis-

similarity between pairs of sites. The Partitioning Around Medoids (PAM) algorithm and silhouette method was then used to map the Gower's distance matrix into a specified set of clusters [7]. PAM identifies medoids, which are the points within each cluster whose sum of dissimilarities to all other sites in the same cluster is minimal, and so, in this case, can be identified as the most representative flux tower sites of each cluster.

Table S1: Site-level variables and definitions used to calculate Gower's distances during cluster analysis.

Variable	Data Type	Definition
ET	Numerical	Mean daily ET rate (mm d^{-1})
T _{air}	Numerical	Mean air temperature at site ($^{\circ}\text{C}$)
SW _{rad}	Numerical	Mean daily incoming short-wave radiation (W m^{-2})
VPD	Numerical	Mean daily vapour pressure deficit (hPa)
Site _{age}	Numerical	Mean forest stand age at site (years)
Wspeed	Numerical	Mean daily wind speed (m s^{-1})
P _{sum_14D}	Numerical	Mean two-week cumulative sum of precipitation prior to the day of measurement (mm)
Lai _{500m}	Numerical	Mean daily leaf area index ($\text{m}^2 \text{m}^{-2}$)
Cover _{type}	Categorical	IGBP cover type classification of site
Climate _{zone}	Categorical	Main Köppen climate group of site

S1.3 Random Forest Modelling and Evaluation Metrics

Random Forest models were trained and tested with `RandomForestRegressor()` from `scikit-learn` [29] using the different methods outlined below. Due to time constraints, a single set of optimised parameter values was calculated once using `GroupKFold()` and `Optuna` [30] to minimise root mean square error (RMSE) on fold predictions. These values were used for all RF models throughout the study (see Table S2).

In all cases, regression analysis metrics of R^2 (coefficient of determination), RMSE, and mean absolute error (MAE) were calculated from model predictions of ET using their respective `scikit-learn` functions alongside observed values of ET at each site. R^2 aims to explain the proportion of variance captured by the model, while RMSE emphasises the size of errors against observed values. MAE was also used due to its simplicity and ability to compare to RMSE for a greater understanding of error structure. While no consensus has been reached on a unified metric to evaluate ML models and predictive accuracy, a model with the the highest values of R^2 and lowest values of RMSE is deemed the most accurate here [31, 3].

Table S2: Parameters selected for use in all RF training setups in this study. Calculated to minimise RMSE on trials using `GroupKFold()` and `Optuna`.

RF parameter	Value
n_estimators	350
max_depth	17
min_samples_leaf	19
max_features	0.35

Table S3: All Features and definitions included in different RF model setups during training and inference.

Variable	Definition
ET	Mean daily ET rate (mm d^{-1})
Site_age	Mean forest stand age at site (years)
T_air	Mean air temperature at site ($^{\circ}\text{C}$)
SW_rad	Mean daily incoming short-wave radiation (W m^{-2})
VPD	Mean daily vapour pressure deficit (hPa)
Wspeed	Mean daily wind speed (m s^{-1})
P_sum_14D	Mean two-week cumulative sum of precipitation prior to the day of measurement (mm)
Lai_500m	Mean daily leaf area index ($\text{m}^2 \text{m}^{-2}$)
sin_D0Y	Cyclical day of year representation alongside cos_D0Y
cos_D0Y	Cyclical day of year representation alongside sin_D0Y

S1.3.1 Model Sensitivity Testing

To assess the sensitivity of the RF models to different features, a forward feature selection assessment was carried out. Non-target (ET) numerical features were added sequentially and the RF model was re-trained and tested by forming daily ET predictions at the medoid sites with each new feature added. Features were added in the following order: T_air, SW_rad, VPD, Wspeed, P_sum_14D, Lai_500m, sin_D0Y, cos_D0Y, Site_Age. Training sites consisted of all sites apart from medoids (see [Table S4](#)), with predictions formed and testing carried out at medoid sites.

S1.3.2 Random Sampling of Training Datasets

To investigate how predictive performance is impacted by including site age as a feature, two separate RF architectures were formed: 'Mod' ([Equation S1](#)) and 'Mod + Age' ([Equation S2](#)).

$$ET_i = RF \left(T_{air_i}, SW_{rad_i}, VPD_i, W_{speed_i}, \sum_{i=13}^i P, LAI_i, \sin_DOY_i, \cos_DOY_i \right) \quad (S1)$$

$$ET_i = RF \left(T_{air_i}, SW_{rad_i}, VPD_i, W_{speed_i}, \sum_{i=13}^i P, LAI_i, \sin_DOY_i, \cos_DOY_i, Site_Age_i \right) \quad (S2)$$

where the target ET_i is the estimated ET rate on the i th day (mm/day); $RF(\cdot)$ is the Random Forest model; the features T_{air} , SW_{rad} , VPD , W_{speed} , $\sum_{i=13}^i P$, LAI , $Site_Age$ are all listed with descriptions in [Table S3](#).

To further assess how the overall ecological similarity of training dataset composition to the test site may impact predictive accuracy of daily ET, a method to subset sites from a larger site pool before RF training was also developed allowing Gower's distances of training sites to be varied and potential effects investigated (script 21_analysis_1.2.py), with details on how this analysis runs below.

Sites were first separated as follows: 6 medoid sites for testing and the remaining 44 sites as the full training pool. Following this, 20 sites were chosen at random (NumPy.random module) from the training pool, and then two models - 'Mod' and 'Mod + Age' - are trained on this subset of sites. After training, both models then predict on the six held out medoid sites, and the mean Gower's distance and mean age distance from the full training subset to each test site is calculated and stored alongside prediction performance metrics. This paired process is then repeated a further 99 times, each with a new NumPy seed to avoid using the same subset, to produce 100 runs using this method.

this is where i need to explain gower's distance mean plot etc.

S1.3.3 Two Site Case Study: US-UMd and DE-Lnf

- details on Leave One Group Out (LOGO) approach - random shuffle details - heatmap and Δ Predicted ET details

S1.4 Site Metadata

Table S4: Metadata summary of FLUXNET sites used in this study (** denotes a cluster medoid)

FLUXNET ID	Days of Data	Country	Climate Zone	Cover Type	Site Age (years)	Mean ET 90D Peak (mm d ⁻¹)
BE-Bra	3590	Belgium	Temperate	MF	92	2.02
BE-Vie	4466	Belgium	Temperate	MF	107	2.15
CA-Gro	3639	Canada	Continental	MF	84	2.45
CA-Man	2290	Canada	Continental	ENF	173	1.90
CA-NS1	1141	Canada	Continental	ENF	157	1.59
CA-NS2	960	Canada	Continental	ENF	76	1.75
CA-NS3	1141	Canada	Continental	ENF	42	1.70
CA-NS5	1137	Canada	Continental	ENF	26	1.88
CA-NS6	1141	Canada	Continental	OSH	17	1.35
CA-NS7	1141	Canada	Continental	OSH	8	1.73
CA-Oas	3029	Canada	Continental	DBF	91	2.82
CA-Obs***	3029	Canada	Continental	ENF	122	1.95
CA-Qfo	2687	Canada	Continental	ENF	106	1.75
CA-TP3***	4550	Canada	Continental	ENF	44	2.45
CA-TPD	1082	Canada	Continental	DBF	100	2.31
CH-Lae	3782	Switzerland	Temperate	MF	190	4.31
CZ-BK1	3920	Czechia	Continental	ENF	37	1.73
DE-Hai	2720	Germany	Temperate	DBF	260	2.68
DE-Lnf***	2670	Germany	Temperate	DBF	122	2.60
DE-Obe	2463	Germany	Temperate	ENF	79	2.21
DE-Tha	4502	Germany	Temperate	ENF	131	2.14
DK-Sor	4482	Denmark	Temperate	DBF	98	2.87
FI-Hyy	4466	Finland	Continental	ENF	59	2.15
FI-Sod	4454	Finland	Continental	ENF	91	1.87
FR-Fon	3586	France	Temperate	DBF	166	3.64
FR-LBr	2346	France	Temperate	ENF	44	2.76
FR-Pue	4550	France	Temperate	EBF	73	1.86
IT-Col	2908	Italy	Temperate	DBF	195	2.58
IT-Cp2	1082	Italy	Temperate	EBF	65	2.38
IT-Cpz	2359	Italy	Temperate	EBF	65	2.02
IT-PT1	897	Italy	Temperate	DBF	15	3.51
IT-Ren	3253	Italy	Continental	ENF	199	3.53
IT-SR2	716	Italy	Temperate	ENF	65	2.21
IT-SRo***	3819	Italy	Temperate	ENF	63	2.51
NL-Loo	4498	Netherlands	Temperate	ENF	119	2.64
RU-Fyo	4458	Russia	Continental	ENF	247	2.51
US-Blo	1901	U.S.A	Temperate	ENF	21	4.04
US-GBT	898	U.S.A	Continental	ENF	181	2.66
US-GLE	3638	U.S.A	Continental	ENF	190	2.36
US-Ha1	3262	U.S.A	Continental	DBF	112	2.40
US-MMS	4538	U.S.A	Temperate	DBF	105	3.23
US-Me3***	2012	U.S.A	Temperate	ENF	23	1.91
US-NR1	4550	U.S.A	Continental	ENF	121	2.71
US-Oho	3615	U.S.A	Continental	DBF	55	4.02
US-PFa	4522	U.S.A	Continental	MF	164	2.35
US-Prr	1355	U.S.A	Continental	ENF	101	1.56
US-Syv	2723	U.S.A	Continental	MF	307	2.52
US-UMB	4538	U.S.A	Continental	DBF	102	3.16
US-UMd***	2715	U.S.A	Continental	DBF	94	2.91
US-WCr	2931	U.S.A	Continental	DBF	106	2.56